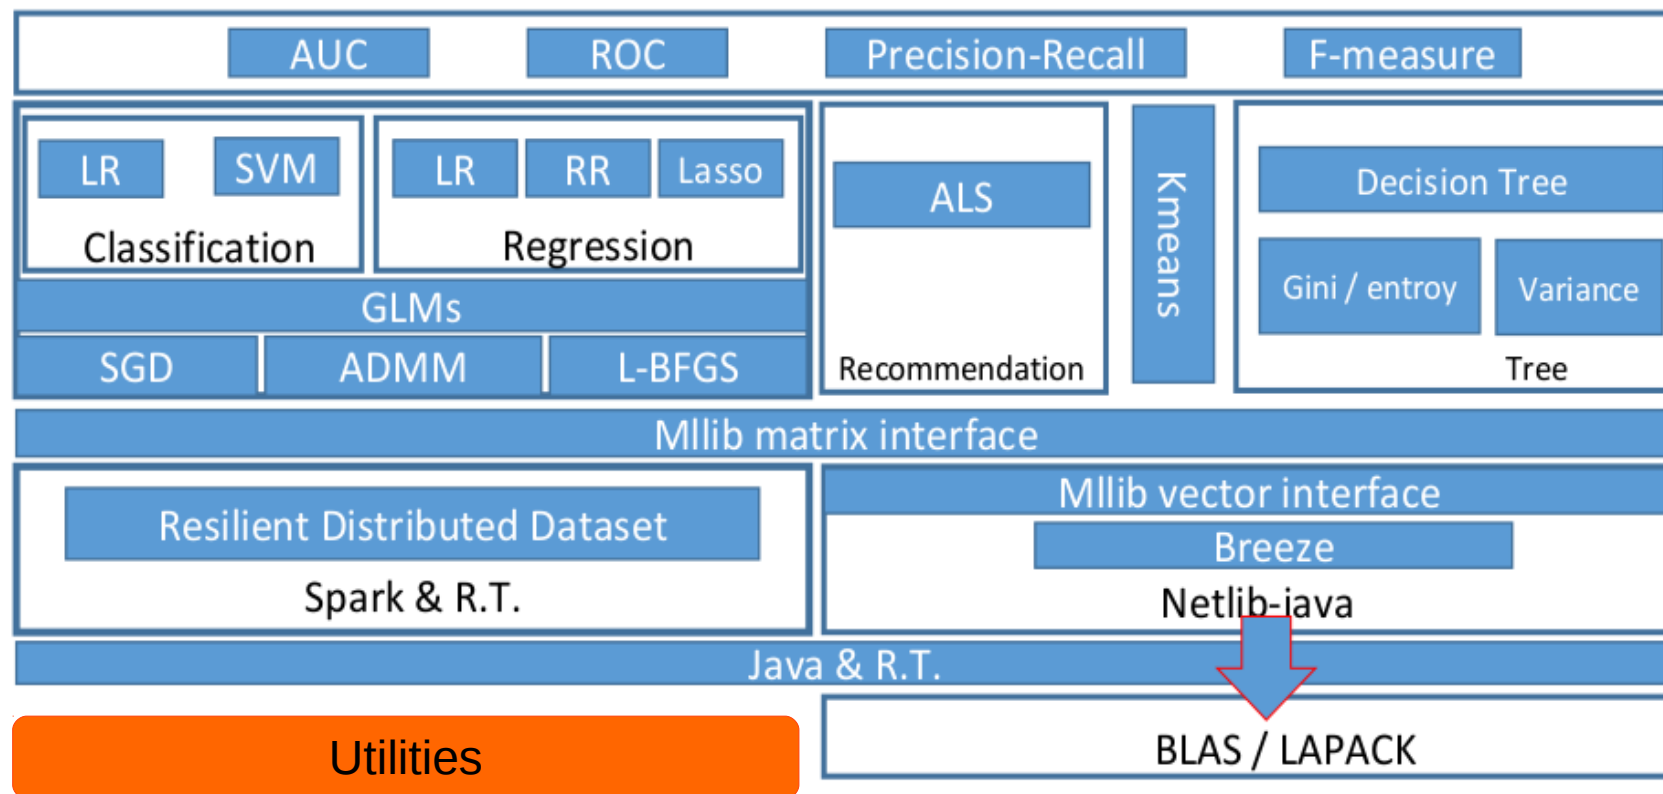


第七课：GraphX 入门

【声明】 本视频和幻灯片为炼数成金网络课程的教学资料，所有资料只能在课程内使用，不得在课程以外范围散播，违者将可能被追究法律和经济责任。

课程详情访问炼数成金培训网站

<http://edu.dataguru.cn>



- MLib 架构
- K-Means
- 协同过滤

- **GraphX 简介**
 - 图的定义
 - GraphX 图处理
 - GraphX 架构
 - GraphX 操作
- **实例演示**
 - 图例演示
 - PageRank 演示


什么是图

- 图是表示物件与物件之间的关系的方法，是图论的基本研究对象。
- 图是由一些小圆点（称为顶点）和连结这些圆点的直线或曲线（称为边）组成的。

基本术语 [编辑]

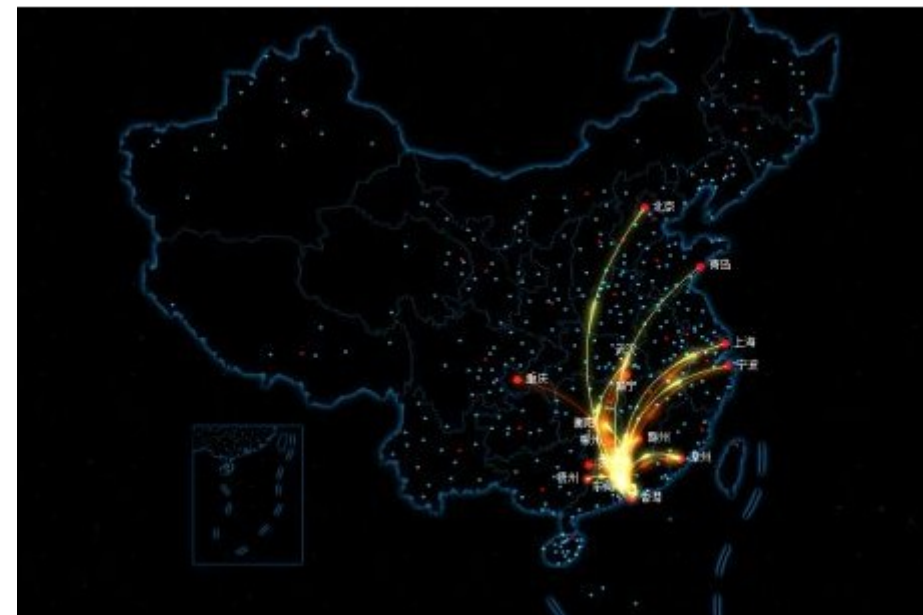

参见：图论术语

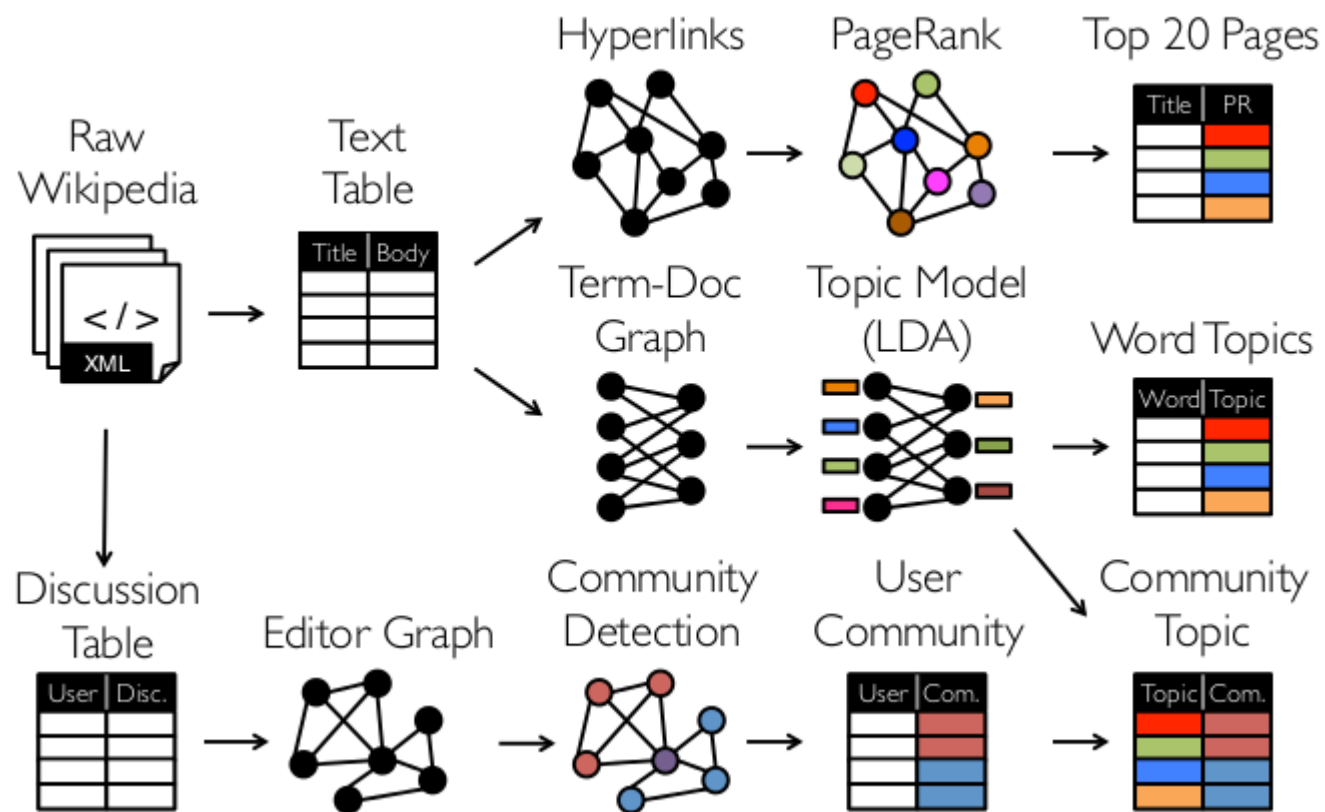
- **阶 (Order)**：图 G 中顶点集 V 的大小称作图 G 的阶。
 - **子图 (Sub-Graph)**：图 G' 称作图 G 的子图如果 $V(G') \subseteq V(G)$ 以及 $E(G') \subseteq E(G)$ 。
 - **生成子图 (Spanning Sub-Graph)**：指满足条件 $V(G') = V(G)$ 的 G 的子图 G 。
 - **度 (Degree)** 是一个顶点的度是指与该顶点相关联的总边数，顶点 v 的度记作 $d(v)$ 。度和边有如下关系：
$$\sum_{v \in V} d(v) = 2|E|$$
 - **出度 (Out-degree) 和入度 (In-degree)**：对有向图而言，顶点的度还可分为出度和入度。一个顶点的出度为 d_o ，是指有 d_o 条边以该顶点为起点，或说与该点关联的出边共有 d_o 条。入度的概念也类似。
 - **邻接矩阵**
 - **自环 (Loop)**：若一条边的两个顶点相同，则此边称作自环。
 - **路径 (Path)**：从顶点 u 到顶点 v 的一条路径是指一个序列 $v_0, e_1, v_1, e_2, v_2, \dots, e_k, v_k$ ， e_i 的起点终点为 v_{i-1} 及 v_i ； k 称作路径的长度； $v_0 = u$ ，称为路径的起点； $v_k = v$ ，称为路径的终点。如果 $u = v$ ，称该路径是闭的，反之则称为开的；如果 v_1, \dots, v_k 两两不等，则称之为**简单路径 (Simple path)**，注意， $u = v$ 是允许的)。
 - **行迹 (Trace)**：如果路径 $P(u, v)$ 中边各不相同，则该路径称为 u 到 v 的一条行迹。
 - **轨道 (Track)**：即简单路径。
 - 闭的行迹称作**回路 (Circuit)**，闭的轨道称作**圈 (Cycle)**。
- (现存文献中的命名法并无统一标准。比如在另一种定义中，walk对应上述的path，path对应上述的track，trail对应上述的trace。)
- **距离 (Distance)**：从顶点 u 出发到顶点 v 的最短路径若存在，则此路径的长度称作从 u 到 v 的距离。若从 u 到 v 根本不存在路径，则记该距离为无穷(∞)。
 - **距离矩阵**
 - **桥 (Bridge)**：若去掉一条边，便会使得整个图不连通，该边称为桥。


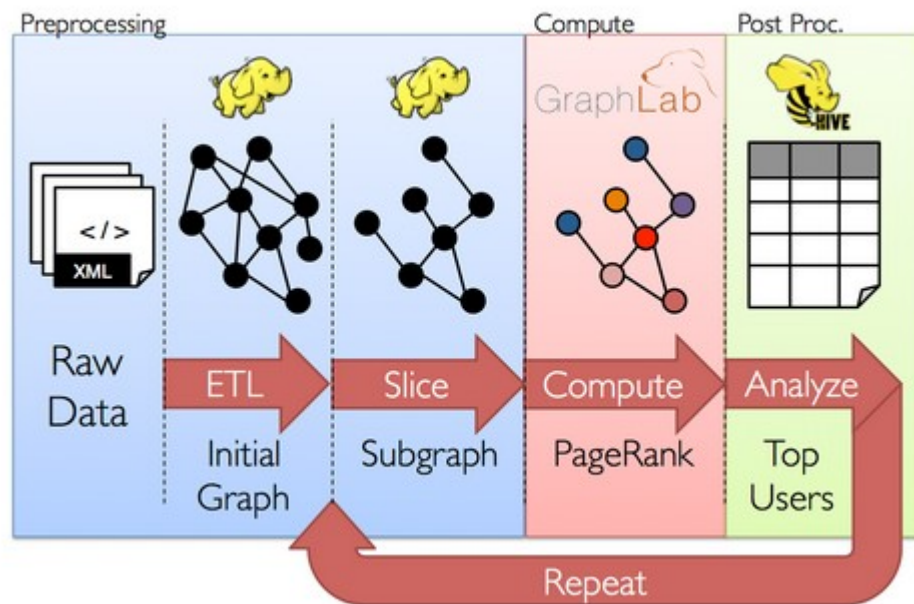


图的应用

- PageRank
- 社交网络
- 网络交易
- 微博
- 交通状况监控
- 监控电商卖家刷评分
-








New API

*Blurs the distinction between
Tables and Graphs*




New System

*Combines Data-Parallel
Graph-Parallel Systems*



Property Graph



Vertex Property Table

Id	Property (V)
Rxin	(Stu., Berk.)
Jegonzal	(PstDoc, Berk.)
Franklin	(Prof., Berk)
Istoica	(Prof., Berk)

Edge Property Table

SrcId	DstId	Property (E)
rxin	jegonzal	Friend
franklin	rxin	Advisor
istoica	franklin	Coworker
franklin	jegonzal	PI

```
class VertexRDD[@specialized VD: ClassTag](  
  val partitionsRDD: RDD[ShippableVertexPartition[VD]],  
  val targetStorageLevel: StorageLevel = StorageLevel.MEMORY_ONLY  
  extends RDD[(VertexId, VD)](partitionsRDD.context, List(new OneToOneDependency(partitionsRDD))) {
```

```
/**  
 * A 64-bit vertex identifier that uniquely identifies a vertex within a graph. It does not need  
 * to follow any ordering or any constraints other than uniqueness.  
 */  
type VertexId = Long
```



顶点



```
class EdgeRDD[@specialized ED: ClassTag, VD: ClassTag] (  
  val partitionsRDD: RDD[(PartitionID, EdgePartition[ED, VD])],  
  val targetStorageLevel: StorageLevel = StorageLevel.MEMORY_ONLY  
  extends RDD[Edge[ED]](partitionsRDD.context, List(new OneToOneDependency(partitionsRDD))) {
```


```
case class Edge[@specialized(Char, Int, Boolean, Byte, Long, Float, Double) ED] (  
  var srcId: VertexId = 0,  
  var dstId: VertexId = 0,  
  var attr: ED = null.asInstanceOf[ED])  
  extends Serializable {
```

```
class EdgeDirection private (private val name: String) extends Serializable {  
  /**  
   * Reverse the direction of an edge. An in becomes out,  
   * out becomes in and both and either remain the same.  
   */  
  def reverse: EdgeDirection = this match {  
    case EdgeDirection.In => EdgeDirection.Out  
    case EdgeDirection.Out => EdgeDirection.In  
    case EdgeDirection.Either => EdgeDirection.Either  
    case EdgeDirection.Both => EdgeDirection.Both  
  }  
}
```


边

```
class EdgeTriplet[VD, ED] extends Edge[ED] {  
  /**...*/  
  var srcAttr: VD = _ // nullValue[VD]  
  
  /**...*/  
  var dstAttr: VD = _ // nullValue[VD]  
  
  /**...*/  
  protected[spark] def set(other: Edge[ED]): EdgeTriplet[VD,ED] = {  
    srcId = other.srcId  
    dstId = other.dstId  
    attr = other.attr  
    this  
  }  
}
```

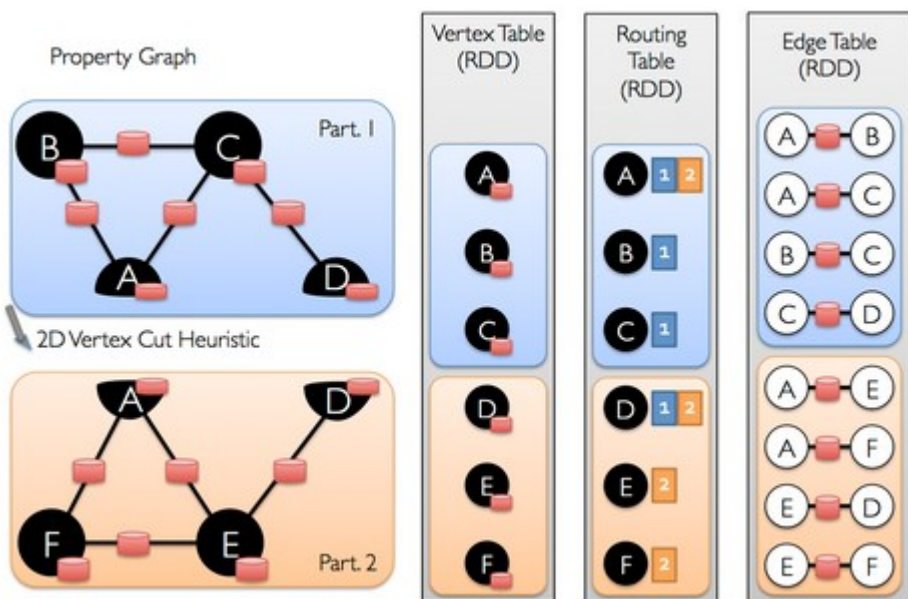




Edge Cut



Vertex Cut



```
Structure
▼ PartitionStrategy
  getPartition(VertexId, VertexId, PartitionId)
  from(String): PartitionStrategy
  EdgePartition2D
  EdgePartition1D
  RandomVertexCut
  CanonicalRandomVertexCut
```

GraphX 图处理


```
Structure
▼ Graph
  v vertices
  v edges
  v triplets
  f persist(StorageLevel): Graph[VD, ED]
  f cache(): Graph[VD, ED]
  f unpersistVertices(Boolean): Graph[VD, ED]
  f partitionBy(PartitionStrategy): Graph[VD, ED]
  f partitionBy(PartitionStrategy, Int): Graph[VD, ED]
  f mapVertices[VD2: ClassTag]((VertexId, VD) => VD2): Graph[VD2, ED]
  f mapEdges[ED2: ClassTag](Edge[ED] => ED2): Graph[VD, ED2]
  f mapEdges[ED2: ClassTag]((PartitionID, Edge[ED]) => ED2): Graph[VD, ED2]
  f mapTriplets[ED2: ClassTag](EdgeTriplet[VD, ED] => ED2): Graph[VD, ED2]
  f mapTriplets[ED2: ClassTag]((PartitionID, EdgeTriplet[VD, ED]) => ED2): Graph[VD, ED2]
  f reverse: Graph[VD, ED]
  f subgraph(EdgeTriplet[VD, ED] => Boolean): Graph[VD, ED]
  f mask[VD2: ClassTag, ED2: ClassTag](Collection[EdgeTriplet[VD, ED]]): Graph[VD2, ED2]
  f groupEdges((ED, ED) => ED): Graph[VD, ED]
  f mapReduceTriplets[A: ClassTag](EdgeTriplet[VD, ED] => A): Graph[VD, A]
  f outerJoinVertices[U: ClassTag, VD2: ClassTag](Graph[VD2, ED], U => VD2): Graph[VD, ED]
  v ops
  f this()

▼ Graph
  f fromEdgeTuples[VD: ClassTag](RDD[(VertexId, VD)]): Graph[VD, ED]
  f fromEdges[VD: ClassTag, ED: ClassTag](RDD[(VertexId, Edge[ED])]): Graph[VD, ED]
  f apply[VD: ClassTag, ED: ClassTag](RDD[(VertexId, Edge[ED])]): Graph[VD, ED]
  f graphToGraphOps[VD: ClassTag, ED: ClassTag]: GraphOps[VD, ED]
```

```
Structure
▼ GraphImpl
  f this()
  v edges
  v triplets
  f persist(StorageLevel): Graph[VD, ED]
  f cache(): Graph[VD, ED]
  f unpersistVertices(Boolean): Graph[VD, ED]
  f partitionBy(PartitionStrategy): Graph[VD, ED]
  f partitionBy(PartitionStrategy, Int): Graph[VD, ED]
  f reverse: Graph[VD, ED]
  f mapVertices[VD2: ClassTag]((VertexId, VD) => VD2): Graph[VD2, ED]
  f mapEdges[ED2: ClassTag]((PartitionID, Edge[ED]) => ED2): Graph[VD, ED2]
  f mapTriplets[ED2: ClassTag]((PartitionID, EdgeTriplet[VD, ED]) => ED2): Graph[VD, ED2]
  f subgraph(EdgeTriplet[VD, ED] => Boolean): Graph[VD, ED]
  f mask[VD2: ClassTag, ED2: ClassTag](Collection[EdgeTriplet[VD, ED]]): Graph[VD2, ED2]
  f groupEdges((ED, ED) => ED): Graph[VD, ED]
  f mapReduceTriplets[A: ClassTag](EdgeTriplet[VD, ED] => A): Graph[VD, A]
  f outerJoinVertices[U: ClassTag, VD2: ClassTag](Graph[VD2, ED], U => VD2): Graph[VD, ED]
  f accessesVertexAttr[AnyRef, String](VertexId, String): AnyRef
  f this(VertexRDD[VD], ReplicatedVertexIds): GraphImpl

▼ GraphImpl
  f apply[VD: ClassTag, ED: ClassTag](RDD[(VertexId, Edge[ED])]): Graph[VD, ED]
  f fromEdgePartitions[VD: ClassTag, ED: ClassTag](Iterable[EdgePartition]): Graph[VD, ED]
  f apply[VD: ClassTag, ED: ClassTag](RDD[(VertexId, Edge[ED])]): Graph[VD, ED]
  f apply[VD: ClassTag, ED: ClassTag](VertexRDD[VD], EdgeRDD[ED]): Graph[VD, ED]
  f fromExistingRDDs[VD: ClassTag, ED: ClassTag](VertexRDD[VD], EdgeRDD[ED]): Graph[VD, ED]
  f fromEdgeRDD[VD: ClassTag, ED: ClassTag](VertexRDD[VD], EdgeRDD[ED]): Graph[VD, ED]
```

```
▼ GraphOps
  v numEdges
  v numVertices
  v inDegrees
  v outDegrees
  v degrees
  f degreesRDD(EdgeDirection): VertexRDD[Double]
  f collectNeighborIds(EdgeDirection): VertexCollection
  f collectNeighbors(EdgeDirection): VertexCollection
  f collectEdges(EdgeDirection): VertexCollection
  f joinVertices[U: ClassTag](RDD[(VertexId, U)]): Graph[VD, ED]
  f filter[VD2: ClassTag, ED2: ClassTag](VD => Boolean, Edge[ED] => Boolean): Graph[VD2, ED2]
  f pickRandomVertex(): VertexId
  f pregel[A: ClassTag](A, Int, EdgeDirection, (A, EdgeDirection) => A): Graph[VD, ED]
  f pageRank(Double, Double): Graph[Double, ED]
  f staticPageRank(Int, Double): Graph[Double, ED]
  f connectedComponents(): Graph[VertexId, ED]
  f triangleCount(): Graph[Int, ED]
  f stronglyConnectedComponents(Int): Graph[VertexId, ED]
  f this(Graph[VD, ED])
```



GraphX 图处理


```
object GraphLoader extends Logging {  
  /**...*/  
  def edgeListFile(  
    sc: SparkContext,  
    path: String,  
    canonicalOrientation: Boolean = false,  
    minEdgePartitions: Int = 1,  
    edgeStorageLevel: StorageLevel = StorageLevel.MEMORY_ONLY,  
    vertexStorageLevel: StorageLevel = StorageLevel.MEMORY_ONLY)  
    : Graph[Int, Int] =  
  {  
    val startTime = System.currentTimeMillis  
  
    // Parse the edge data table directly into edge partitions  
    val lines = sc.textFile(path, minEdgePartitions).coalesce(minEdgePartitions)  
    val edges = lines.mapPartitionsWithIndex { (pid, iter) =>  
      val builder = new EdgePartitionBuilder[Int, Int]  
      iter.foreach { line =>  
        if (!line.isEmpty && line(0) != '#') {  
          val lineArray = line.split("\\s+")  
          if (lineArray.length < 2) {  
            logWarning("Invalid line: " + line)  
          }  
          val srcId = lineArray(0).toLong  
          val dstId = lineArray(1).toLong  
          if (canonicalOrientation && srcId > dstId) {  
            builder.add(dstId, srcId, 1)  
          } else {  
            builder.add(srcId, dstId, 1)  
          }  
        }  
      }  
    }  
  }
```

```
/**  
 * Loads a graph from an edge list formatted file where each line contains two integers  
 * id and a target id. Skips lines that begin with `#`.  
 *  
 * If desired the edges can be automatically oriented in the positive  
 * direction (source Id < target Id) by setting `canonicalOrientation` to  
 * true.  
 *  
 * @example Loads a file in the following format:  
 * {{{  
 * # Comment Line  
 * # Source Id <\t> Target Id  
 * 1 -5  
 * 1 2  
 * 2 7  
 * 1 8  
 * }}}  
 *  
 * @param sc SparkContext  
 * @param path the path to the file (e.g., /home/data/file or hdfs://file)  
 * @param canonicalOrientation whether to orient edges in the positive  
 * direction  
 * @param minEdgePartitions the number of partitions for the edge RDD  
 * @param edgeStorageLevel the desired storage level for the edge partitions  
 * @param vertexStorageLevel the desired storage level for the vertex partitions  
 */
```

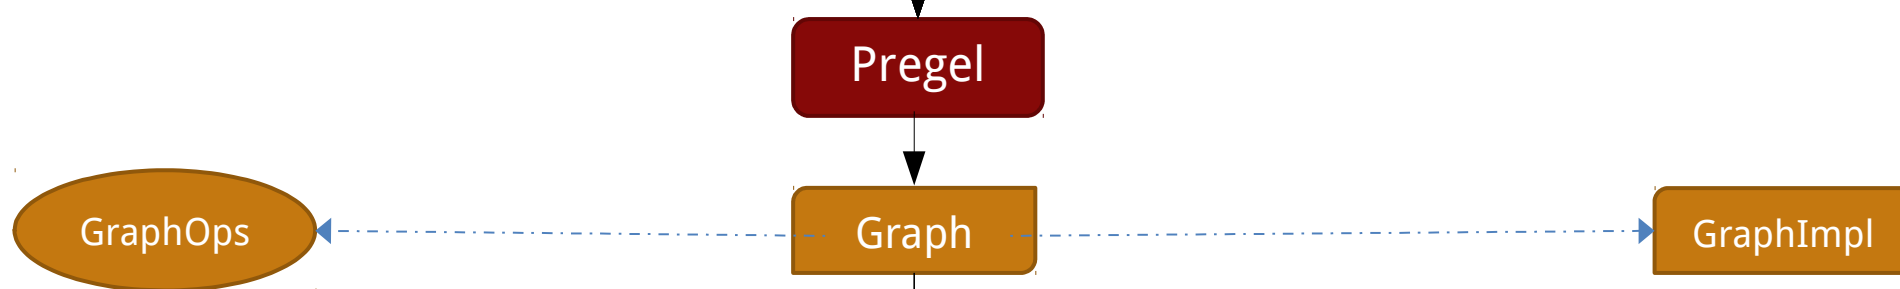


graph


算法



模型



实现



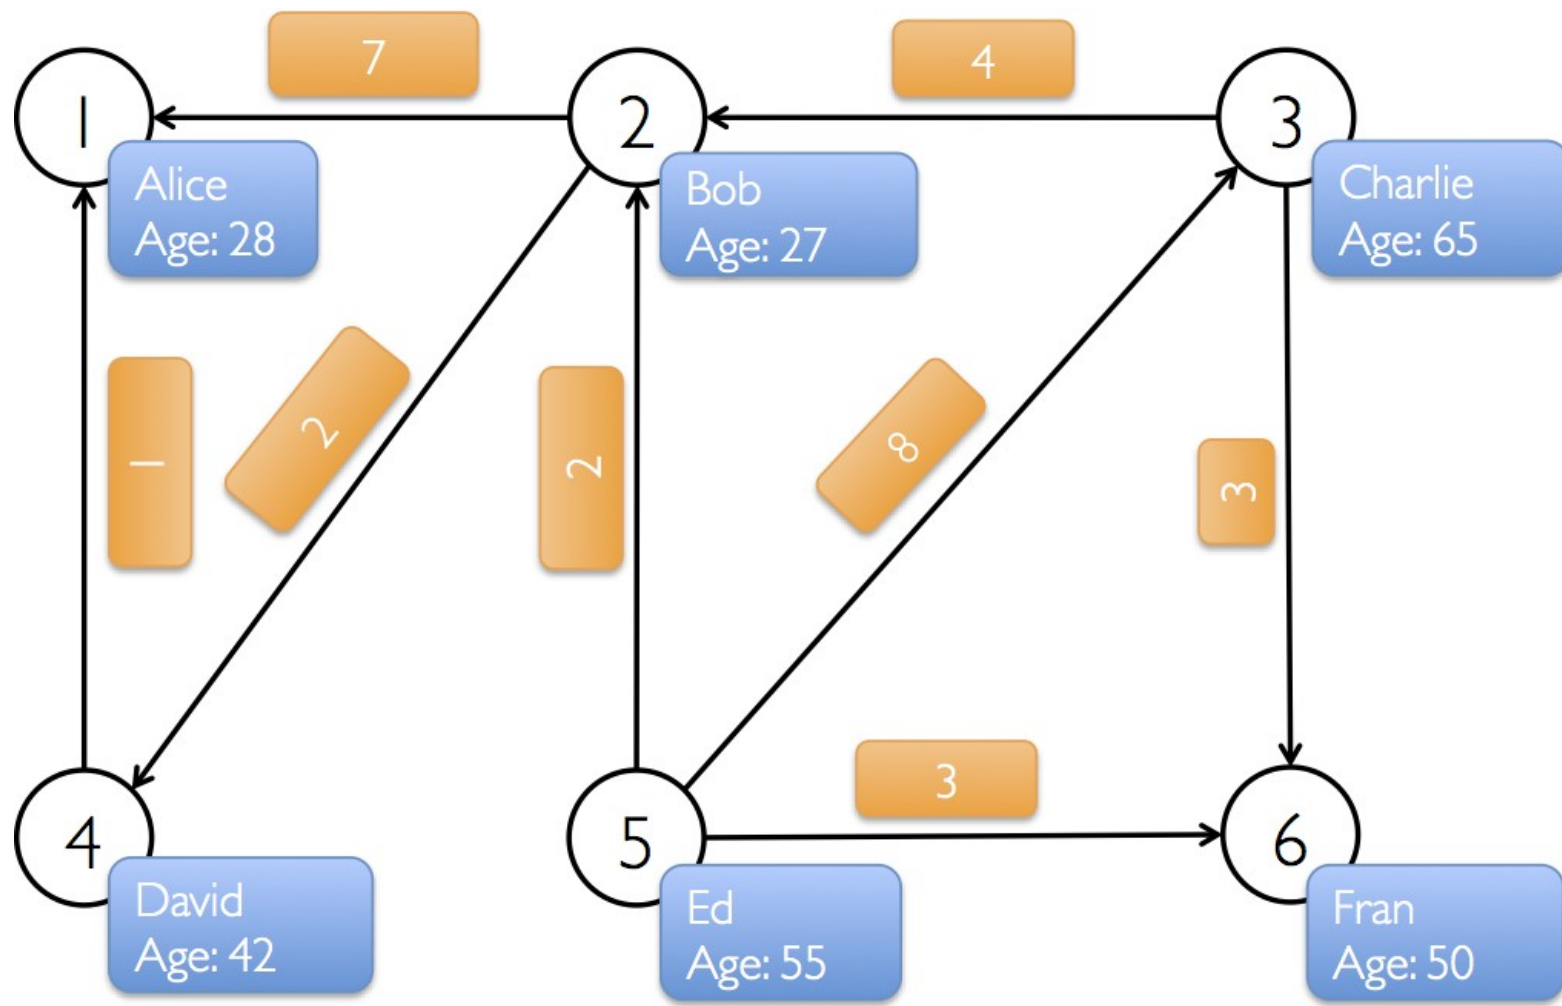


```
/** Summary of the functionality in the property graph */
class Graph[VD, ED] {
  // Information about the Graph =====
  val numEdges: Long
  val numVertices: Long
  val inDegrees: VertexRDD[Int]
  val outDegrees: VertexRDD[Int]
  val degrees: VertexRDD[Int]
  // Views of the graph as collections =====
  val vertices: VertexRDD[VD]
  val edges: EdgeRDD[ED, VD]
  val triplets: RDD[EdgeTriplet[VD, ED]]
  // Functions for caching graphs =====
  def persist(newLevel: StorageLevel = StorageLevel.MEMORY_ONLY): Graph[VD, ED]
  def cache(): Graph[VD, ED]
  def unpersistVertices(blocking: Boolean = true): Graph[VD, ED]
  // Change the partitioning heuristic =====
  def partitionBy(partitionStrategy: PartitionStrategy): Graph[VD, ED]
  // Transform vertex and edge attributes =====
  def mapVertices[VD2](map: (VertexID, VD) => VD2): Graph[VD2, ED]
  def mapEdges[ED2](map: Edge[ED] => ED2): Graph[VD, ED2]
  def mapEdges[ED2](map: (PartitionID, Iterator[Edge[ED]]) => Iterator[ED2]): Graph[VD, ED2]
  def mapTriplets[ED2](map: EdgeTriplet[VD, ED] => ED2): Graph[VD, ED2]
  def mapTriplets[ED2](map: (PartitionID, Iterator[EdgeTriplet[VD, ED]]) => Iterator[ED2])
    : Graph[VD, ED2]
```

```
// Modify the graph structure =====
def reverse: Graph[VD, ED]
def subgraph(
  epred: EdgeTriplet[VD, ED] => Boolean = (x => true),
  vpred: (VertexID, VD) => Boolean = ((v, d) => true))
  : Graph[VD, ED]
def mask[VD2, ED2](other: Graph[VD2, ED2]): Graph[VD, ED]
def groupEdges(merge: (ED, ED) => ED): Graph[VD, ED]
// Join RDDs with the graph =====
def joinVertices[U](table: RDD[(VertexID, U)])(mapFunc: (VertexID, VD, U) => VD): Graph[VD, ED]
def outerJoinVertices[U, VD2](other: RDD[(VertexID, U)])
  (mapFunc: (VertexID, VD, Option[U]) => VD2)
  : Graph[VD2, ED]
// Aggregate information about adjacent triplets =====
def collectNeighborIds(edgeDirection: EdgeDirection): VertexRDD[Array[VertexID]]
def collectNeighbors(edgeDirection: EdgeDirection): VertexRDD[Array[(VertexID, VD)]]
def mapReduceTriplets[A: ClassTag](
  mapFunc: EdgeTriplet[VD, ED] => Iterator[(VertexID, A)],
  reduceFunc: (A, A) => A,
  activeSetOpt: Option[(VertexRDD[_, EdgeDirection]) = None]
  : VertexRDD[A]
```

- GraphX 简介
 - 图的定义
 - GraphX 图处理
 - GraphX 架构
 - GraphX 操作

- 实例演示
 - 图例演示
 - PageRank 演示




```
override def mapReduceTriplets[A: ClassTag](  
  mapFunc: EdgeTriplet[VD, ED] => Iterator[(VertexId, A)],  
  reduceFunc: (A, A) => A,  
  activeSetOpt: Option[(VertexRDD[_], EdgeDirection)] = None): VertexRDD[A] = {
```

Map-Reduce for each vertex

$\text{mapF}((A \leftarrow B) \Rightarrow A_1$

$\text{mapF}((A \leftarrow C) \Rightarrow A_2$

$\text{reduceF}(A_1, A_2) \Rightarrow A$



- PageRank, 即网页排名, 又称网页级别、Google 左侧排名或佩奇排名。
 - 是 Google 创始人拉里·佩奇和谢尔盖·布林于 1997 年构建早期的搜索系统原型时提出的链接分析算法。
 - 目前很多重要的链接分析算法都是在 PageRank 算法基础上衍生出来的。
- PageRank 是 Google 用于用来标识网页的等级 / 重要性的一种方法, 是 Google 用来衡量一个网站的好坏的唯一标准。
- 在揉合了诸如 Title 标识和 Keywords 标识等所有其它因素之后, Google 通过 PageRank 来调整结果, 使那些更具“等级 / 重要性”的网页在搜索结果中令网站排名获得提升, 从而提高搜索结果的相关性和质量。
- 其级别从 0 到 10 级, 10 级为满分。PR 值越高说明该网页越受欢迎 (越重要)。



这是Google最核心的算法，用于给每个网页价值评分，是Google “在垃圾中找黄金” 的关键算法，这个算法成就了今天的Google


PageRank vector q is defined as $q = Gq$

where $G = \alpha S + (1 - \alpha) \frac{1}{n} U$

- ❑ S is the destination-by-source stochastic matrix,
- ❑ U is all one matrix.
- ❑ n is the number of nodes
- ❑ α is the weight between 0 and 1 (e.g., 0.85)

Algorithm: Iterative powering for finding the first eigen-vector

$$q^{next} = Gq^{cur}$$



$$G = \begin{bmatrix} 0 & 0 & 0 & 0 \\ 1/3 & 0 & 0 & 1 \\ 1/3 & 1/2 & 0 & 0 \\ 1/3 & 1/2 & 1 & 0 \end{bmatrix}$$

- GraphX 构成
- GraphX 操作
- 下周运维和优化
 - 运维（参数配置、HistoryServer、Jobserver、ganglia 集成）
 - 优化（序化、压缩。。。）

See You Next



Thanks

FAQ 时间